

class09

Andres Rivero

10/27/2021

```
fna.data <- "WisconsinCancer.csv"
```

```
wisc.df <- read.csv(fna.data, row.names=1)  
head(wisc.df)
```

```
##      diagnosis radius_mean texture_mean perimeter_mean area_mean  
## 842302      M      17.99      10.38      122.80      1001.0  
## 842517      M      20.57      17.77      132.90      1326.0  
## 84300903     M      19.69      21.25      130.00      1203.0  
## 84348301      M      11.42      20.38       77.58       386.1  
## 84358402      M      20.29      14.34      135.10      1297.0  
## 843786      M      12.45      15.70       82.57       477.1  
##      smoothness_mean compactness_mean concavity_mean concave.points_mean  
## 842302      0.11840      0.27760      0.3001      0.14710  
## 842517      0.08474      0.07864      0.0869      0.07017  
## 84300903     0.10960      0.15990      0.1974      0.12790  
## 84348301     0.14250      0.28390      0.2414      0.10520  
## 84358402     0.10030      0.13280      0.1980      0.10430  
## 843786      0.12780      0.17000      0.1578      0.08089  
##      symmetry_mean fractal_dimension_mean radius_se texture_se perimeter_se  
## 842302      0.2419      0.07871      1.0950      0.9053      8.589  
## 842517      0.1812      0.05667      0.5435      0.7339      3.398  
## 84300903     0.2069      0.05999      0.7456      0.7869      4.585  
## 84348301     0.2597      0.09744      0.4956      1.1560      3.445  
## 84358402     0.1809      0.05883      0.7572      0.7813      5.438  
## 843786      0.2087      0.07613      0.3345      0.8902      2.217  
##      area_se smoothness_se compactness_se concavity_se concave.points_se  
## 842302     153.40      0.006399      0.04904      0.05373      0.01587  
## 842517      74.08      0.005225      0.01308      0.01860      0.01340  
## 84300903     94.03      0.006150      0.04006      0.03832      0.02058  
## 84348301     27.23      0.009110      0.07458      0.05661      0.01867  
## 84358402     94.44      0.011490      0.02461      0.05688      0.01885  
## 843786     27.19      0.007510      0.03345      0.03672      0.01137  
##      symmetry_se fractal_dimension_se radius_worst texture_worst  
## 842302     0.03003      0.006193      25.38      17.33  
## 842517     0.01389      0.003532      24.99      23.41  
## 84300903     0.02250      0.004571      23.57      25.53  
## 84348301     0.05963      0.009208      14.91      26.50  
## 84358402     0.01756      0.005115      22.54      16.67  
## 843786     0.02165      0.005082      15.47      23.75  
##      perimeter_worst area_worst smoothness_worst compactness_worst
```

```
wisc.data <- wisc.df[,-1]
diagnosis <- factor(wisc.df$diagnosis, levels = c("B", "M"))
print(diagnosis)
```

Q1: How many observations are in this dataset?

```
## [1] 569 30
```

Q2: How many of the observations have a malignant diagnosis?

```
table(diagnosis)
```

```
## diagnosis
##      B      M
## 357 212
```

There is 212 malignant observations

****Q3: How many variables/features in the data are suffixed with `_mean`?****

```
length(grep("_mean", colnames(wisc.data)))
```

```
## [1] 10
```

10 columns have the term “`_mean`” in them

Next: PCA

Check column means and standard deviations

```
colMeans(wisc.data)
```

##	radius_mean	texture_mean	perimeter_mean
##	1.412729e+01	1.928965e+01	9.196903e+01
##	area_mean	smoothness_mean	compactness_mean
##	6.548891e+02	9.636028e-02	1.043410e-01
##	concavity_mean	concave.points_mean	symmetry_mean
##	8.879932e-02	4.891915e-02	1.811619e-01
##	fractal_dimension_mean	radius_se	texture_se
##	6.279761e-02	4.051721e-01	1.216853e+00
##	perimeter_se	area_se	smoothness_se
##	2.866059e+00	4.033708e+01	7.040979e-03
##	compactness_se	concavity_se	concave.points_se
##	2.547814e-02	3.189372e-02	1.179614e-02
##	symmetry_se	fractal_dimension_se	radius_worst
##	2.054230e-02	3.794904e-03	1.626919e+01
##	texture_worst	perimeter_worst	area_worst
##	2.567722e+01	1.072612e+02	8.805831e+02
##	smoothness_worst	compactness_worst	concavity_worst
##	1.323686e-01	2.542650e-01	2.721885e-01
##	concave.points_worst	symmetry_worst	fractal_dimension_worst
##	1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data,2,sd)
```

##	radius_mean	texture_mean	perimeter_mean
##	3.524049e+00	4.301036e+00	2.429898e+01
##	area_mean	smoothness_mean	compactness_mean
##	3.519141e+02	1.406413e-02	5.281276e-02
##	concavity_mean	concave.points_mean	symmetry_mean
##	7.971981e-02	3.880284e-02	2.741428e-02
##	fractal_dimension_mean	radius_se	texture_se

```
##          7.060363e-03          2.773127e-01          5.516484e-01
##          perimeter_se          area_se          smoothness_se
##          2.021855e+00          4.549101e+01          3.002518e-03
##          compactness_se          concavity_se          concave.points_se
##          1.790818e-02          3.018606e-02          6.170285e-03
##          symmetry_se          fractal_dimension_se          radius_worst
##          8.266372e-03          2.646071e-03          4.833242e+00
##          texture_worst          perimeter_worst          area_worst
##          6.146258e+00          3.360254e+01          5.693570e+02
##          smoothness_worst          compactness_worst          concavity_worst
##          2.283243e-02          1.573365e-01          2.086243e-01
##          concave.points_worst          symmetry_worst          fractal_dimension_worst
##          6.573234e-02          6.186747e-02          1.806127e-02
```

```
wisc.pr <- prcomp(wisc.data, scale. = TRUE)
```

```
summary(wisc.pr)
```

```
## Importance of components:
```


```
##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation    3.6444  2.3857  1.67867  1.40735  1.28403  1.09880  0.82172
## Proportion of Variance 0.4427  0.1897  0.09393  0.06602  0.05496  0.04025  0.02251
## Cumulative Proportion 0.4427  0.6324  0.72636  0.79239  0.84734  0.88759  0.91010
##          PC8      PC9      PC10     PC11     PC12     PC13     PC14
## Standard deviation    0.69037  0.6457  0.59219  0.5421  0.51104  0.49128  0.39624
## Proportion of Variance 0.01589  0.0139  0.01169  0.0098  0.00871  0.00805  0.00523
## Cumulative Proportion 0.92598  0.9399  0.95157  0.9614  0.97007  0.97812  0.98335
##          PC15     PC16     PC17     PC18     PC19     PC20     PC21
## Standard deviation    0.30681  0.28260  0.24372  0.22939  0.22244  0.17652  0.1731
## Proportion of Variance 0.00314  0.00266  0.00198  0.00175  0.00165  0.00104  0.0010
## Cumulative Proportion 0.98649  0.98915  0.99113  0.99288  0.99453  0.99557  0.9966
##          PC22     PC23     PC24     PC25     PC26     PC27     PC28
## Standard deviation    0.16565  0.15602  0.1344  0.12442  0.09043  0.08307  0.03987
## Proportion of Variance 0.00091  0.00081  0.0006  0.00052  0.00027  0.00023  0.00005
## Cumulative Proportion 0.99749  0.99830  0.9989  0.99942  0.99969  0.99992  0.99997
##          PC29     PC30
## Standard deviation    0.02736  0.01153
## Proportion of Variance 0.00002  0.00000
## Cumulative Proportion 1.00000  1.00000
```

Q4: From your results, what proportion of the original variance is captured by the first principal components (PC1)? 44.27%

Q5: How many principal components (PCs) are required to describe at least 70% of the original variance in the data? 3

Q6: How many principal components (PCs) are required to describe at least 90% of the original variance in the data? 7

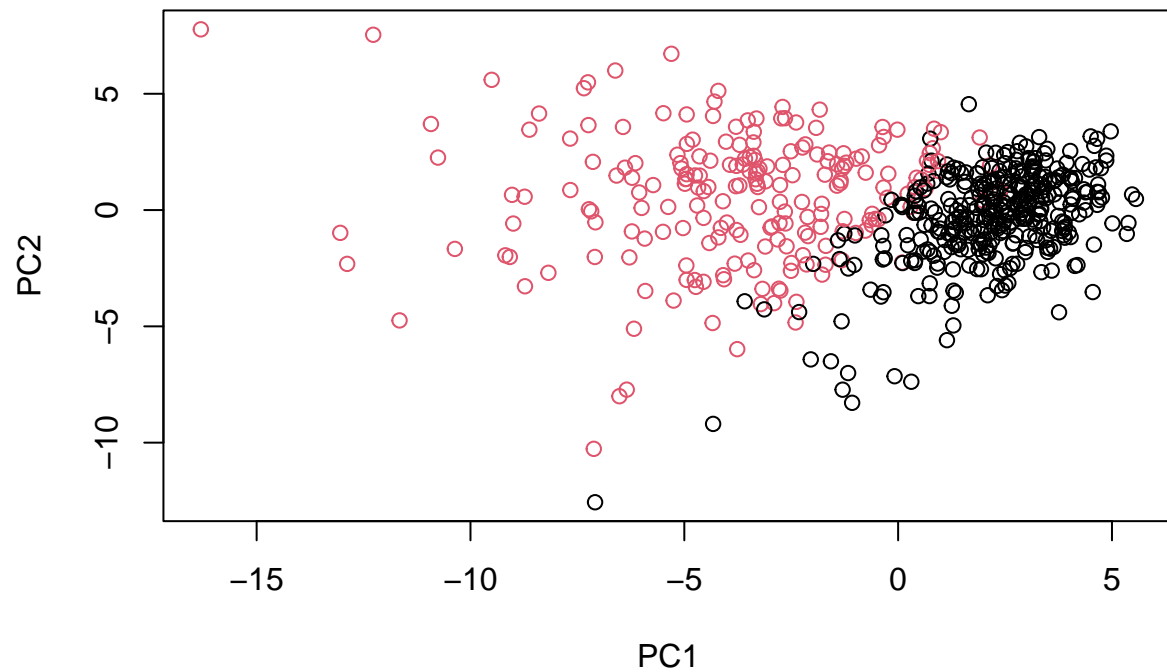
```
biplot(wisc.pr)
```



Q7: What stands out to you about this plot? Is it easy or difficult to understand? Why?


Can't understand anything on the plot, impossible to read

```
plot( wisc.pr$x , col = diagnosis ,
      xlab = "PC1", ylab = "PC2")
```



Q8: Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
plot( wisc.pr$x[, c(1,3)] , col = diagnosis ,  
      xlab = "PC1", ylab = "PC3")
```




There seems to be two different cluster of cells, red dots(or malignant) and black dots (or benign)

Creating ggplot for our data

```
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis
```

```
library(ggplot2)
ggplot(df) + aes(PC1, PC2, col=diagnosis) +
  geom_point()
```

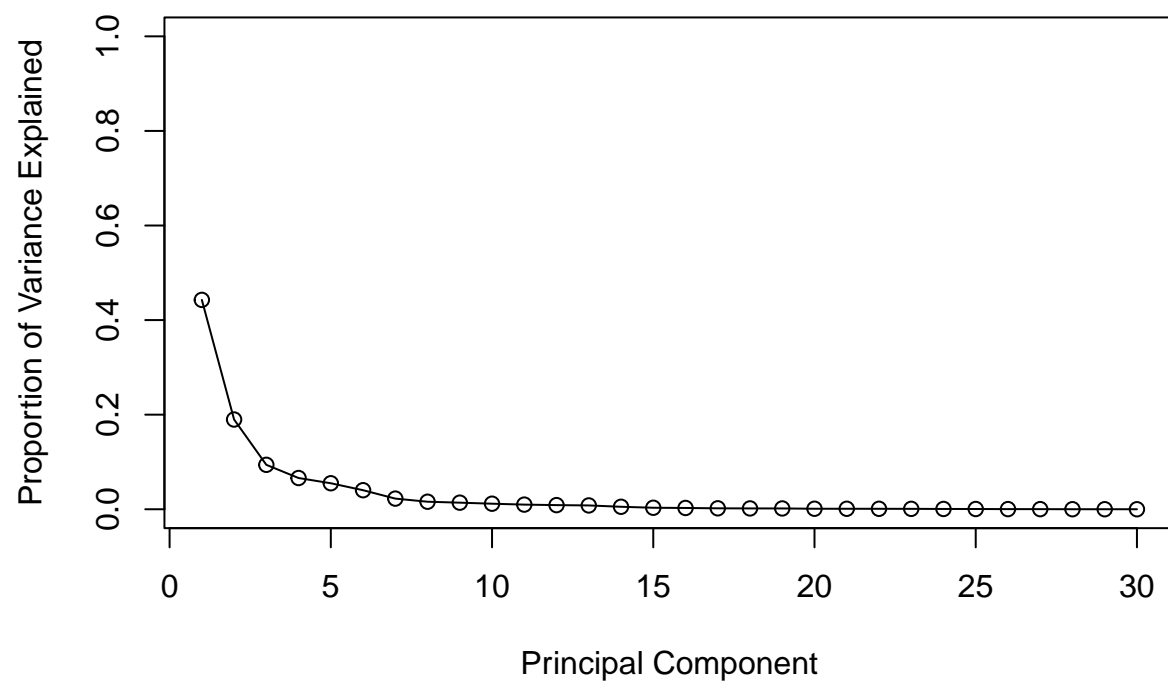


Calculating the variance now


```
pr.var <- wisc.pr$sdev^2  
head(pr.var)
```

```
## [1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

```
pve <- pr.var/(sum(pr.var))  
plot(pve, xlab = "Principal Component",  
      ylab = "Proportion of Variance Explained",  
      ylim = c(0, 1), type = "o")
```

```
barplot(pve, ylab = "Precent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```



Q9: For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

```
wisc.pr$rotation["concave.points_mean",1]
```

```
## [1] -0.2608538
```

Q10: What is the minimum number of principal components required to explain 80% of the variance of the data?

```
summary(wisc.pr)
```

```
## Importance of components:
##               PC1    PC2    PC3    PC4    PC5    PC6    PC7
## Standard deviation  3.6444 2.3857 1.67867 1.40735 1.28403 1.09880 0.82172
## Proportion of Variance 0.4427 0.1897 0.09393 0.06602 0.05496 0.04025 0.02251
## Cumulative Proportion 0.4427 0.6324 0.72636 0.79239 0.84734 0.88759 0.91010
##               PC8    PC9    PC10   PC11   PC12   PC13   PC14
## Standard deviation  0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624
## Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
## Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
##               PC15   PC16   PC17   PC18   PC19   PC20   PC21
## Standard deviation  0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
```


```
##          PC22    PC23    PC24    PC25    PC26    PC27    PC28
## Standard deviation  0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
##          PC29    PC30
## Standard deviation  0.02736 0.01153
## Proportion of Variance 0.00002 0.00000
## Cumulative Proportion 1.00000 1.00000
```

We need The first 5 PCs to get to at least 80%

Q11: Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters? *Hierarchical clustering*

```
data.scaled <- scale(wisc.data)
data.dist <- dist(data.scaled)
wisc.hclust <- hclust(data.dist)
plot(wisc.hclust) +
  abline(h=19, col="red", lty=2)
```

Cluster Dendrogram



```
data.dist
hclust (*, "complete")
```

```
## integer(0)
```

Height 19 gives us 4 clusters.

```
wisc.hclust.clusters <- cutree(wisc.hclust, h=19)
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##              1 12 165
##              2  2   5
##              3 343  40
##              4  0   2
```

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

```
wisc.hclust.clusters <- cutree(wisc.hclust, h=10)
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##              1   0  24
##              2   0  53
##              3   0   3
##              4   0  11
##              5  12  46
##              6 100  33
##              7   0  14
##              8 197   6
##              9   0   5
##             10  19   0
##             11  11   0
##             12   1   0
##             13  10   0
##             14   0   2
##             15   0   6
##             16   1   0
##             17   0   6
##             18   4   0
##             19   0   2
##             20   0   1
##             21   2   0
```

```
wisc.hclust.clusters <- cutree(wisc.hclust, h=20)
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##              1 12 165
##              2  2   5
##              3 343  40
##              4  0   2
```


```
wisc.hclust.clusters <- cutree(wisc.hclust, h=24)
table(wisc.hclust.clusters, diagnosis)
```

```
##           diagnosis
## wisc.hclust.clusters  B  M
##           1 357 210
##           2   0   2
```

It seems that height 19-20 gives us the best clusters.

Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:4]), method="ward.D2")
plot(wisc.pr.hclust)
abline(h=80, col="red")
```



```
dist(wisc.pr$x[, 1:4])
hclust (*, "ward.D2")
```

Q15: How well does the newly created model with four clusters separate out the two diagnoses?

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
## grps
##    1    2
## 171 398
```

Comparing it to the expert M and B vector

```
table(diagnosis)
```

```
## diagnosis
##   B   M
## 357 212
```

We can do a cross-table by giving the table() function two inputs. (Called confusion table)

```
table(grps, diagnosis)
```

```
##      diagnosis
## grps   B   M
##    1   6 165
##    2 351  47
```

Accuracy, essentially how many did we get correct?

```
(165+351)/nrow(wisc.data)
```

```
## [1] 0.9068541
```

Q17: Which of your analysis procedures resulted in a clustering model with the best specificity?
How about sensitivity?

Sensitivity

```
165/(165+47)
```

```
## [1] 0.7783019
```

Specificity

```
351/(351+47)
```


```
## [1] 0.8819095
```

```
url <- "new_samples.csv"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

```
##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
## [1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
## [2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
##          PC8      PC9      PC10     PC11     PC12     PC13     PC14
## [1,] -0.2307350  0.1029569 -0.9272861  0.3411457  0.375921  0.1610764  1.187882
## [2,] -0.3307423  0.5281896 -0.4855301  0.7173233 -1.185917  0.5893856  0.303029
##          PC15     PC16     PC17     PC18     PC19     PC20
```

```
## [1,] 0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
## [2,] 0.1299153 0.1448061 -0.40509706 0.06565549 0.25591230 -0.4289500
##      PC21      PC22      PC23      PC24      PC25      PC26
## [1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
## [2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
##      PC27      PC28      PC29      PC30
## [1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
## [2,] -0.001134152 0.09638361 0.002795349 -0.019015820
```

```
plot(wisc.pr$x[,1:2], col=diagnosis)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], labels=c(1,2), col="white")
```



Q18. Which of these new patients should we prioritize for follow up based on your results?
 Patient number 2